

One-page note — Kill the Quote Spreadsheet

What this is: A short note for anyone reviewing the live demo. **Live demo:** <https://kill-the-quote-spreadsheet-lac.vercel.app/> **Code:** <https://github.com/Anirudhmen94/Aerchainfinal>

What we are doing (in one paragraph)

Procurement buyers still spend days copying vendor quotes into Excel. They send an RFx to five suppliers and get back messy replies — Excel that ignores the template, a PDF with a discount buried in a footnote, a Word paragraph, a phone photo of a rate card, or a one-line email. Then a manager asks one “what if” question and another afternoon disappears. **This prototype ends that week.** A buyer drafts the RFx with an AI co-pilot, vendors reply however they like, the system reads every format into one side-by-side comparison (same lines, same units, same currency), and the buyer asks questions in plain English through to an award decision they can defend. **Demo size:** corrugated packaging for a Chakan snacks plant · 5 vendors · 30 line items · 8 quality questions.

How the product works (walk the tabs)

Tab	What happens
Draft	Buyer describes the need in plain language → AI drafts ~30 lines + quality questionnaire + 5 vendors.
Outbox	Preview cover emails → “send” writes stub files (no real email server). Later: award notices and manager alerts land here too.
Inbox	Seed or upload the five ugly reply formats → Parse → structured quotes.
Compare	One INR matrix for all vendors. Click a cell to see the original source (snippet + file). Unclear values are marked uncertain/missing — not hidden.
Ask	Chat in normal English over the comparison (e.g. cheapest split among vendors who passed quality). Answers come from real extracted data, not canned scripts.
Award	Pick a vendor per line → save. If you want a failing/incomplete vendor on a line, the system notifies a manager via Outbox (there is no fake Approve button).
Audit	Log of saves, notices, and overrides — so decisions are reviewable.

Trust rules we built for: every price should link back to evidence; math (FX, unit conversion, cheapest-line) runs in code, not in the AI; vendors who fail knockout questions are clearly flagged and kept off the default award list.

What we deliberately left out (and why)

- **Real email (SMTP/IMAP).** Outbox saves stub files only. We prove the messages and workflow without building a mail product.
- **Freeze / Unfreeze on Award.** Removed on purpose. Saving the award + Audit is enough for a demo; a freeze ceremony added ceremony, not trust.
- **Vendor portal or forced reply template.** Vendors stay free to reply in any format; our job is to absorb the mess.
- **Approve / Reject UI for managers on Award.** Managers get a notify stub in Outbox only. The buyer owns the screen; we did not fake a full approval workflow.
- **ERP, payments, purchase orders.** Scope stops at a defensible award decision.
- **Multi-user login / permissions.** Single-buyer demo so the ugly-edge story stays center stage.
- **Hardcoded answers in Ask.** Free questions must run on the extracted data — canned answers would break the assignment.
- **Silently applying footnote discounts into award totals.** Discounts stay visible as caveats; quiet math into “official” totals would destroy trust.
- **One-click “run everything” on Home.** The reviewer walks the real path tab by tab.

Where we think the hard problem actually is

Reading documents is getting easier. The hard product problem is **matching messy vendor lines to the right RFx rows** when data is partial (“same as last year”, only 27 of 30 lines quoted, incomplete quality answers). The honest product answer is clear flags, buyer override, and an audit trail — not a silent best guess. Next worth building: a short “confirm these mappings” link back to the vendor so nobody retypes Excel again.